

Dashboard metrics

Revised MVP priorities

Based on: Dashboard Metrics Catalogue.pdf, 8 pages. Preserve Impact, Voice Ops and Insights; prioritize real demo evidence and a clear owner action.

Scope decision: the new catalogue proposes estimated net saving as the hero. Your earlier requirement to defer assumptions still applies, so bookings and covers remain the hero. The priorities below are recommendations for this MVP, not a reproduction of the source's P0 labels.

P0 / Impact - four tiles, ranked

- 1. Bookings and covers (1.12).** Count bookings and sum party sizes. Shows a tangible business outcome. Include comparable previous-period counts; do not equate bookings with attended visits or revenue.
- 2. Calls handled (1.3).** Count distinct calls and show previous-period comparison. Gives workload context and makes small demo samples visible.
- 3. Containment (1.8).** Calls without a Callback Request / calls. Shows handoff demand; explain "No callback requested", rather than claiming successful resolution.
- 4. Calls outside opening hours (1.14).** Count and share starting outside configured hours. Demonstrates availability beyond business hours, without claiming those calls would have been missed.

P0 / Voice Ops - five tiles, ranked

- 1. Task success (2.1).** Success / (success + failure), excluding unknown grades. Provides explicit task-completion evidence; display grading coverage.
- 2. Containment (2.2).** Reuse the exact Impact definition, label and band. Lets the technical view stand alone without conflicting numbers.
- 3. LLM TTFB p50/p90 (2.3).** One tile, nearest-rank percentiles over agent turns. Exposes model response speed; explicitly distinguish it from end-to-end caller wait.
- 4. Error rate (2.7).** Calls with configured failed statuses / calls. Exposes technical failures, with missing-status coverage visible.
- 5. Planned vs forced handoffs (2.10 + 2.11).** One two-part tile: callback requests in each class / calls. Distinguishes policy escalation from agent limitations. Include Unclassified; omit assumption-based health bands.

P0 / Insights - three panels, ranked

- 1. Unanswered questions (3.1).** Evidence-ranked review queue with source calls, last-asked time and persistent Approve/Dismiss. Turns gaps into an owner decision.
- 2. Open callbacks and oldest wait (3.4).** Current backlog across all dates, oldest first. Makes waiting customers actionable. Use a validated warning threshold only.
- 3. Handoff reasons (3.2).** Count requests by reason with planned/forced tags and Unclassified. Explains which issues the owner should investigate.

Global P0: freshness banner (3.3), separate from tile counts. Display last refresh and the configured stale warning.

What ships, what waits

Make P0 real before adding more tiles

Build Impact outcomes and Insights follow-through, then complete Voice Ops provider wiring. The catalogue reports that task success, LLM TTFB and status are not yet saved in app/. Their existing tiles are not evidence of a complete integration. Verify payloads and persist these existing provider fields; if missing, display No data and flag demo readiness as incomplete.

This corrects an ordering conflict in the source: its build list starts with several P1 rows despite saying all P0 comes first. The MVP build order follows the selected P0 hierarchy above.

P1 / Add only after the core flow works

- 1. Booking conversion (1.13).** Distinct Booked calls / calls. Useful funnel context beside bookings; require a validated comparison or approved band, otherwise keep as detail.
- 2. Abandoned rate (2.12).** Abandoned calls / calls. Reveals incomplete journeys; promote when a meaningful comparison or approved band exists.
- 3. Confirmation emails sent (2.14).** Sent bookings / bookings. Verifies follow-through; Sent alone establishes neither delivery nor the 30-second target.
- 4. Busiest times (3.6).** Business-timezone weekday/hour heatmap. Supports staffing decisions once the sample is useful. Do not prioritize a sparse demo heatmap over P0.

Defer assumptions, capture work and secondary detail

Financial estimates: cost pair and net saving (1.1, 1.2, 1.7), avoided hours/labour (1.5, 1.6), sliders (1.16) and revenue scenarios (1.15). The catalogue permits an estimate hero; this brief deliberately defers it under your scope. Connected minutes (1.4) belongs in detail, not a tile.

Unverified or additional instrumentation: assumed traffic baseline (2.8), coherence criterion setup (2.6), measured bill/cap/cost per minute (1.10, 1.11, 2.18), and end-to-end latency capture (2.4). The source promotes the latter to P1; this MVP defers new capture and labels TTFB honestly.

Lower-value duplication: fallback rate (2.9) as optional detail because the gap queue already supports action; average duration (2.13) and review-count analytics (3.5) later. Persisting review decisions remains P0 even though review analytics does not.

Retain source cuts: stage latency, tool breakdown, scripted-test tiles, email attempts, topic capture, revenue/upsell claims and unrelated advanced analytics.

Presentation and trust rules

Default budget: **4 Impact + 5 Voice Ops + 3 Insights**, plus freshness. Keep optional operational tiles within nine. Use one period and business timezone, comparable elapsed periods, visible counts/coverage and No data for empty denominators. Avoid healthy colours when records are incomplete.

Read thresholds from YAML. The source's two-hour callback warning is only an example: do not silently make it an SLA. Without a validated threshold, show age and oldest-first ordering. Do not invent conversion or abandonment targets. Use neutral trend arrows rather than implying every increase is good.

Source trace: catalogue pp. 2-4 (Impact), pp. 4-6 (Voice Ops), pp. 6-7 (Insights), pp. 7-8 (gaps, build order and cuts). Row IDs preserve traceability. Repository behaviour and external benchmark claims have not been independently verified; no benchmark is needed for this selection.

Claude Code prompt / 1 of 2

Rework the existing Forward AI in Business dashboard using Dashboard Metrics Catalogue.pdf as the metric reference. Preserve Impact, Voice Ops and Insights. Apply the hierarchy below, which deliberately defers assumptions and new capture for the hackathon. Inspect actual repository code and data first; catalogue statements are not proof of working integrations. Reuse existing components, styling, queries and authentication. Do not deploy.

A. Inspect Dashboards/, app/, config/dashboard_rules.yaml, business_rules.yaml, liveCalls, liveBookings, liveCallbacks, gap reviews and actual ElevenLabs post-call payloads. Confirm field names, types, timestamps, units and status semantics. Identify which selected metrics can be populated from real data. Never fabricate missing samples or present unknown as zero.

B. Impact: use four P0 tiles in this order: bookings and covers (1.12), calls handled (1.3), containment (1.8), and calls outside opening hours (1.14). Bookings = booking-record count; covers = sum of valid party sizes. Calls = distinct voice sessions. Containment = calls with no Callback Request / calls; use the identical label, calculation and configured band in Voice Ops, with a "No callback requested" explanation. After-hours = calls starting outside configured opening hours / calls; show count and rate. Handle overnight hours, closed days and the business timezone. Do not imply those calls would necessarily have been missed. Add comparable previous-period counts to bookings/covers and calls, with neutral arrows; show No comparison when unavailable and no percentage growth from a zero baseline. Bookings are recorded outcomes, not attended visits or measured revenue.

C. Voice Ops: build five P0 slots ordered by Business, Quality, Operations, Handoffs. Business: task success (2.1) and containment (2.2). Quality: LLM TTFB p50/p90 (2.3), one tile. Operations: error rate (2.7). Handoffs: planned vs forced rates (2.10+2.11), one two-part tile, with Unclassified visible. Persist existing provider fields analysis.call_successful, turn convai_llm_service_ttfb and status through the existing ingestion handler after verifying real payload paths. Preserve webhook security, validate input and make repeated events idempotent. Task success = $S/(S+F)$, excluding unknown; show graded coverage. TTFB uses nearest-rank percentiles over valid turn samples, with verified units and sample count; label "Model time to first byte, not caller-perceived wait". Error rate = calls with status in configured failed_statuses / calls; expose missing-status coverage and suppress healthy bands when evidence is incomplete. Callback rates follow the catalogue: requests in each class / calls; label the denominator and do not silently substitute unique-call rates. Deduplicate request IDs; request rates can exceed 100% if multiple requests belong to a call. Planned reasons: large_party, booking_change, catering, complaint. Forced: unanswered_question, other. Reuse config mappings and retain all other codes as Unclassified. Show handoff counts/rates without assumed healthy targets. If provider data is unavailable, retain No data with a coverage explanation; report that demo readiness is incomplete.

D. Insights: three P0 panels: unanswered questions (3.1), current open callbacks with oldest wait (3.4), handoff-reason counts (3.2). Group questions ignoring case, whitespace and punctuation; sort unreviewed first, then distinct-call evidence count. Show source call IDs, last-asked timestamp, evidence and Approve/Dismiss. Approve collects an answer through the existing review workflow; Dismiss persists a decision. Verify or add minimal review persistence, without auto-publishing invented answers. Callback backlog includes all currently open records across dates and is explicitly labelled Current backlog. Show oldest valid createdAt and wait; use an existing owner-approved age threshold if available. The PDF's two-hour warning is an example, not an established service target. Without a validated threshold, sort oldest first and omit overdue claims. Reason bars show request counts and planned/forced/unclassified tags.

Claude Code prompt / 2 of 2

E. Shared rules: one selected period and configured business timezone; default today. Use explicit, consistent timestamp semantics for bookings versus call cohorts. Use like-for-like elapsed periods for comparisons. Show freshness (3.3) dashboard-wide using asOf and configured stale_after_minutes, catalogue default 60. Distinguish loading, empty, stale, incomplete and error states. Zero denominators yield No data. Show count denominators and coverage. Use text/icons with colour; thresholds come from YAML, never duplicated constants. Do not invent thresholds to satisfy the catalogue's admission test. Clearly label any sample data. Limit the default screen to four Impact tiles, five Voice Ops tiles and three Insights panels plus freshness.

F. After P0 works, P1 may add booking conversion (1.13: distinct Booked calls / calls), abandoned rate (2.12: distinct Abandoned calls / calls), email sent rate (2.14: bookings with Sent / bookings) and busiest-times heatmap (3.6). Require an approved band or valid comparison for conversion and abandonment; otherwise use supporting detail. Sent does not prove delivery or the under-30-second target. Keep optional Voice Ops additions within the catalogue's nine-tile ceiling. Do not let P1 delay missing P0 provider data.

G. Defer net saving and cost/staff comparison (1.7, 1.1, 1.2), avoided time/labour (1.5, 1.6), sliders (1.16), revenue scenarios (1.15), assumed call-volume baseline (2.8), coherence setup (2.6), measured costs/caps (1.10, 1.11, 2.18), new latency capture (2.4), average duration (2.13) and review analytics (3.5). Connected minutes (1.4) may appear only as detail. Knowledge-gap rate (2.9) is optional supporting detail; Insights already exposes its actionable evidence. Preserve reusable code but keep deferred metrics out of the default demo. Retain the catalogue's cuts: stage latency, tool failure breakdown, scripted-test tiles, email attempts, topics and unrelated advanced analytics.

H. Test meaningful cases: empty denominators, unknown grades/statuses, duplicate requests/gaps, multiple bookings per call, percentile ranks and units, overnight hours/timezone boundaries, zero-baseline comparisons, old open callbacks, unclassified reasons and stale data. Verify persisted reviews, ingestion idempotency and the three-tab flow. Run relevant existing checks. Deliver changes, check results, remaining data/configuration gaps and a walkthrough: recorded business outcome -> technical evidence -> owner action.